

Task 1: Train a Regression Model and Tune Hyperparameters

[Github](#) , [LinkedIn](#), email : mannarmannan02@gmail.com

Assigned project Documentation:[Boston model-hparam tuning and CI/CD Implementation](#)

Model selection: XG-Boost

Parallelism: Kubernetes Jobs (Hyper-parameter-tuning)

Introduction:

I implemented a model to predict house prices based on the provided features. And, I developed an end-to-end CI/CD pipeline for this model using GitHub Actions. Below is the file structure and contents for a detailed evaluation.

File 1: train.py

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]

column_names = [
    "CRIM", "ZN", "INDUS", "CHAS", "NOX", "RM", "AGE", "DIS", "RAD", "TAX",
    "PTRATIO", "B", "LSTAT"
]

boston_df = pd.DataFrame(data, columns=column_names)
boston_df['MEDV'] = target

X = boston_df.drop('MEDV', axis=1)
y = boston_df['MEDV']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

#Arguments to be passed from Kubernetes into model container
import argparse
parser = argparse.ArgumentParser()
parser.add_argument("--n_estimators", type=int, required=True)
parser.add_argument("--max_depth", type=int, required=True)
parser.add_argument("--learning_rate", type=float, required=True)
parser.add_argument("--subsample", type=float, required=True)
args = parser.parse_args()

model = XGBRegressor(
    n_estimators=args.n_estimators,
    max_depth=args.max_depth,
    learning_rate=args.learning_rate,
    subsample=args.subsample
)
```

```

model.fit(X_train, y_train)

y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"MSE: {mse}")

model_filename = f"/data/model_{mse}.joblib"
metrics_filename = f"/data/metrics_{mse}.txt"

joblib.dump(model, model_filename)
with open(metrics_filename, "w") as f:
    f.write(str(mse)+"\n")

```

Code Description:

1. Arguments:

The script takes the following hyperparameters as arguments, which are passed via Kubernetes:

- **n_estimators** * **max_depth** * **learning_rate** * **subsample**.

2. Training and Splitting Data:

The Boston housing dataset is split into **80% training and 20% testing**.

3. Model and MSE Scores:

The trained model is saved as **model_{mse}.joblib**. The mean squared error (MSE) is saved as **metrics_{mse}.txt** for further analysis.

File 2: Docker File (Dockerfile.train)

```

FROM python:3.8-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install -r requirements.txt
COPY train.py .
ENTRYPOINT ["python", "train.py"]

```

File 3: Requirements.txt

```

pandas==2.0.3
scikit-learn==1.3.0
xgboost==1.7.6
joblib==1.3.2
flask==2.0.1
numpy==1.21.2

```

File 4: Kubernetes Job Template (job-hyperparameter.yaml)

```

apiVersion: batch/v1
kind: Job
metadata:
  name: trial-{{N_ESTIMATORS}}-{{MAX_DEPTH}}-{{LEARNING_RATE}}-{{SUBSAMPLE}}
  labels:

```

```

app: xgb-hyperparameter-tuning
trial: "true"
spec:
  template:
    metadata:
      labels:
        app: trial
        n_estimators: "{{N_ESTIMATORS}}"
        max_depth: "{{MAX_DEPTH}}"
        learning_rate: "{{LEARNING_RATE}}"
        subsample: "{{SUBSAMPLE}}"
    spec:
      containers:
        - name: xgb-trainer
          image: mannarn/model-train:latest
          args:
            - "--n_estimators={{N_ESTIMATORS}}"
            - "--max_depth={{MAX_DEPTH}}"
            - "--learning_rate={{LEARNING_RATE_RAW}}"
            - "--subsample={{SUBSAMPLE_RAW}}"
          volumeMounts:
            - name: data-volume
              mountPath: /data
      restartPolicy: Never
      volumes:
        - name: data-volume
          hostPath:
            path: /data
            type: DirectoryOrCreate

```

Code Description:

- Kind**
The Job resource in Kubernetes ensures that a pod runs to completion.
- Metadata (Job Name and Labels):**
The job name follows a structured naming convention using hyperparameters:
trial-{{N_ESTIMATORS}}-{{MAX_DEPTH}}-{{LEARNING_RATE}}-{{SUBSAMPLE}}
Example: trial-200-5-0-1-0.6
- Pod Metadata:**
Hyperparameters (**n_estimators**, **max_depth**, **learning_rate**, **subsample**) are assigned as pod labels for easy identification.
- Passing Hyperparameters as Arguments:**
These arguments are passed to the container when it starts, allowing the container to configure the XGBoost model dynamically.

File 5: Script to Launch Trials (generate-jobs.ps1)

```

$n_estimators_list = @(100, 200)
$max_depth_list = @(5, 7)
$learning_rate_list = @(0.01, 0.1)
$subsample_list = @(0.6, 0.8)

```

```
# All combination of h-parameters
foreach ($n in $n_estimators_list) {
    foreach ($depth in $max_depth_list) {
        foreach ($lr in $learning_rate_list) {
            foreach ($sub in $subsample_list) {
                # Sanitize values with dots (e.g., 0.01 → 0-01 for Kubernetes naming)
                $sanitized_lr = "$lr".Replace(".", "-")
                $sanitized_sub = "$sub".Replace(".", "-")

                $yamlContent = (Get-Content job-hyperparameter.yaml -Raw) `
                    -replace '\{\{N_ESTIMATORS\}\}', $n `
                    -replace '\{\{MAX_DEPTH\}\}', $depth `
                    -replace '\{\{LEARNING_RATE\}\}', $sanitized_lr `
                    -replace '\{\{SUBSAMPLE\}\}', $sanitized_sub `
                    -replace '\{\{LEARNING_RATE_RAW\}\}', $lr `
                    -replace '\{\{SUBSAMPLE_RAW\}\}', $sub

                Write-Output "Submitting job: trial-$n-$depth-$sanitized_lr-$sanitized_sub"
                $yamlContent | kubectl apply -f -
            }
        }
    }
}
```

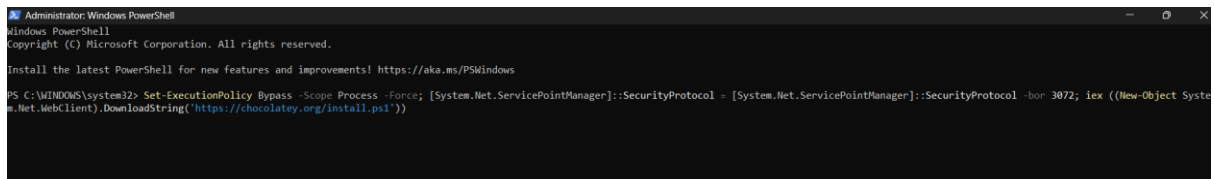
Code Description:

1. **Automated Job Submission:**
This PowerShell script generates Kubernetes job configurations for all possible hyperparameter combinations.
2. **Sanitizing Values for Kubernetes Naming:**
Since Kubernetes does not allow dots (.) in resource names, we replace them with hyphens (-):
`$sanitized_lr = "$lr".Replace(".", "-")`
`$sanitized_sub = "$sub".Replace(".", "-")`
 This ensures compliance with **DNS-1123 Kubernetes naming rules**

STEPS TO RUN Hyperparameter tuning in Local machine

#Install dependencies:

1.Chocoley or Winget to download Mini Kube:



2. Install Mini Kube and Kubernetes CLI tools to create a local development cluster.

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Install the latest PowerShell for new features and improvements! https://aka.ms/PSWindows

PS C:\Users\manna> choco install minikube kubernetes-cli
Chocolatey v2.4.2
Usage of the --trace option is only allowed when running from an elevated session.
3 validations performed. 2 success(es), 1 warning(s), and 0 error(s).

Validation Warnings:
- A pending system reboot request has been detected, however, this is being ignored due to the current Chocolatey configuration. If you want to halt when this occurs, then either set the global feature using:
  choco feature enable --name="exitOnRebootDetected"
  or pass the option --exit-when-reboot-detected.

Chocolatey detected you are not running from an elevated command shell
(cmd/powershell).

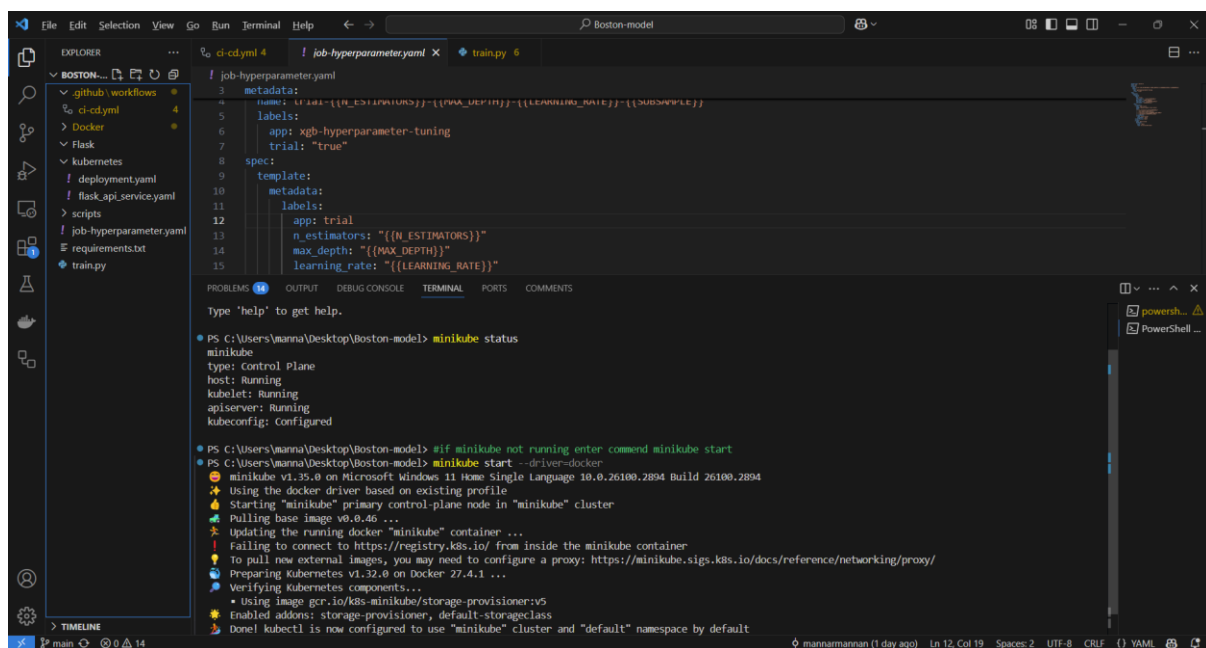
You may experience errors - many functions/packages
require admin rights. Only advanced users should run choco w/out an
elevated shell. When you open the command shell, you should ensure
that you do so with "Run as Administrator" selected. If you are
attempting to use Chocolatey in a non-administrator setting, you
must select a different location other than the default install
location. See
https://docs.chocolatey.org/en-us/choco/setup#non-administrative-install
for details.

Do you want to continue?([Y]es/[N]o): #Press Y to download minikube|
```

3. Start Mini Kube Cluster

>>minikube start --driver=docker

Initializes a local Kubernetes cluster using Docker driver for container orchestration.

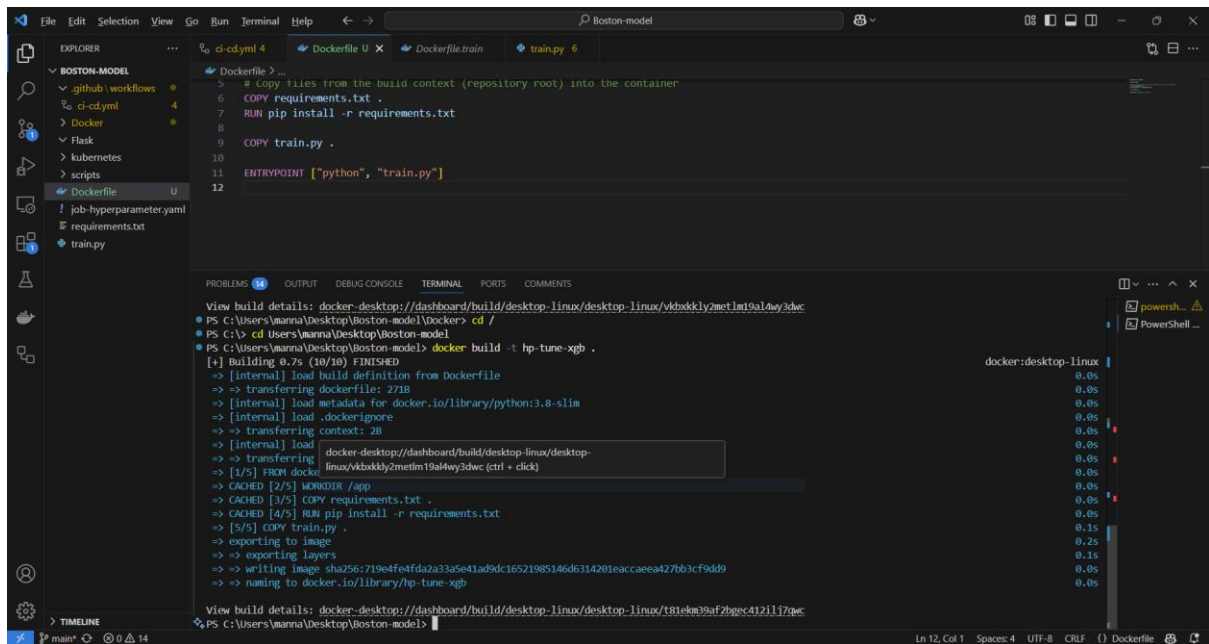


```
File Edit Selection View Go Run Terminal Help
Boston-model
EXPLORER
  BOSTON...
  .github/workflows
  ci-dymil
  Docker
  Flask
  kubernetes
  deployment.yaml
  flask_api_service.yaml
  scripts
  job-hyperparameter.yaml
  requirements.txt
  train.py
  job-hyperparameter.yaml
  metadata:
  3
  4
  5
  6
  7
  8
  9
  10
  11
  12
  13
  14
  15
  labels:
  app: xgb-hyperparameter-tuning
  trial: "true"
  spec:
  template:
  metadata:
  labels:
  app: trial
  n_estimators: "([N_ESTIMATORS])"
  max_depth: "([MAX_DEPTH])"
  learning_rate: "([LEARNING_RATE])"
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS COMMENTS
Type 'help' to get help.
• PS C:\Users\manna\Desktop\Boston-model> minikube status
minikube
type: Control plane
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured
• PS C:\Users\manna\Desktop\Boston-model> #if minikube not running enter command minikube start
• PS C:\Users\manna\Desktop\Boston-model> minikube start --driver=docker
minikube v1.35.0 on Microsoft Windows 11 Home Single Language 10.0.26100.2894 Build 26100.2894
* Using the docker driver based on existing profile
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.46 ...
* Updating the running docker "minikube" container ...
! Failing to connect to https://registry.k8s.io/ from inside the minikube container
! to pull new external images, you may need to configure a proxy: https://minikube.sigs.k8s.io/docs/reference/networking/proxy/
* Preparing Kubernetes V1.32.0 on Docker 27.4.1 ...
* Verifying Kubernetes components...
  * Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

4 Build Training Image

>>docker build -t hp-xgboost .

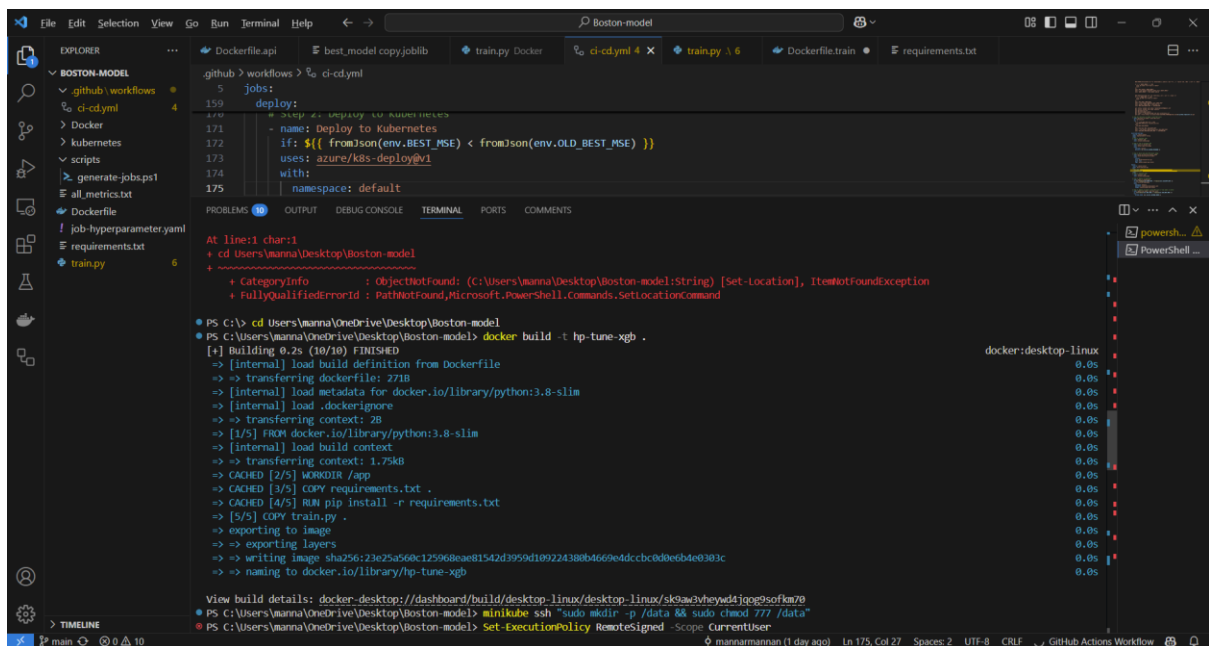
Creates Docker image with training environment and XGBoost dependencies.



5. Create Shared Volume

```
>>minikube ssh "sudo mkdir -p /data && sudo chmod 777 /data"
```

Sets up persistent storage in Minikube VM for model artifacts and metrics.



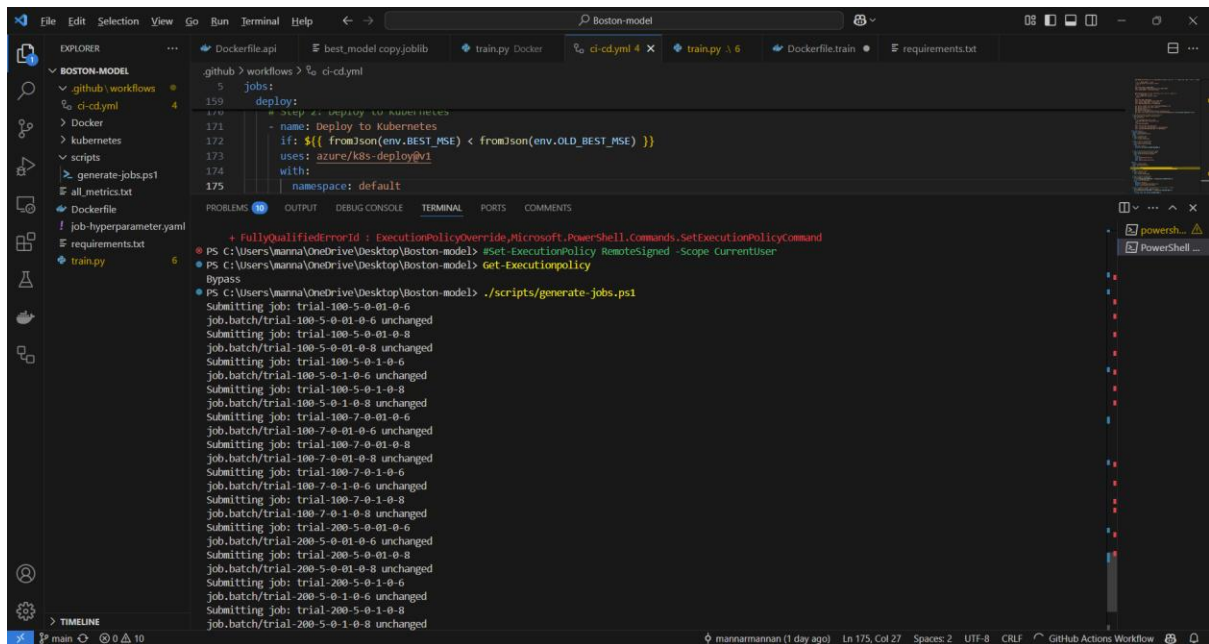
5. Enable PowerShell Script Execution & Launch Parallel Training Jobs

```
>>Set-ExecutionPolicy Bypass -Scope Process
```

Temporarily allows PowerShell script execution for hyperparameter job generation.

```
>>./scripts/generate-jobs.ps1
```

Submits 16 Kubernetes jobs with different hyperparameter combinations for distributed tuning.

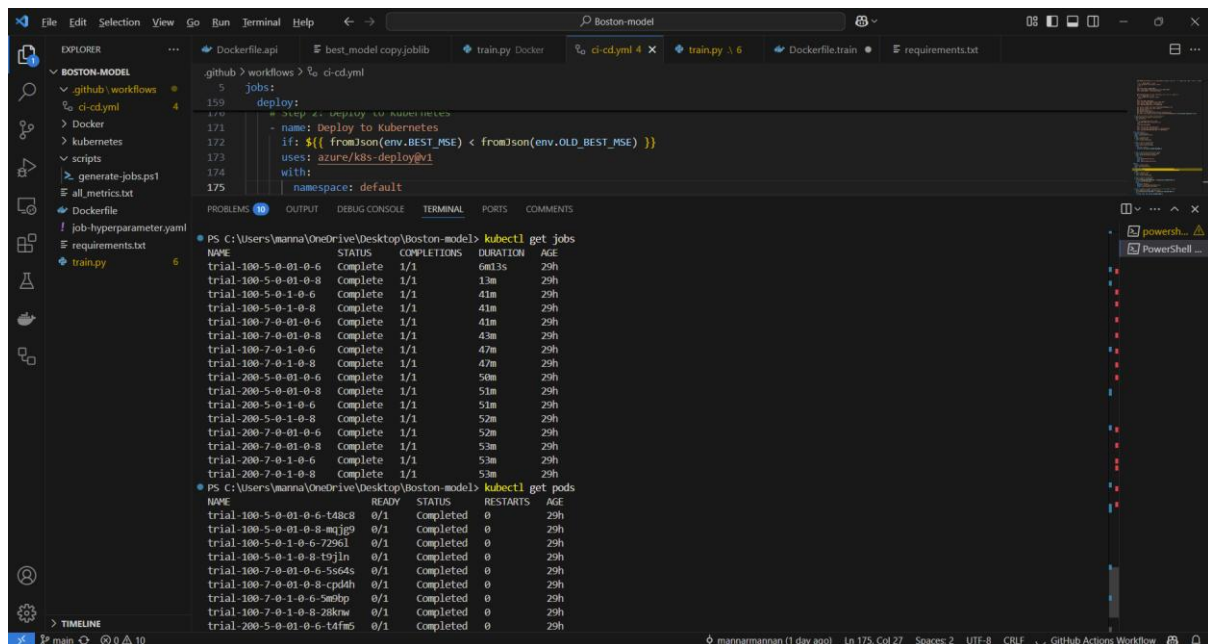


6. Monitor Job Progress

```
>>kubectl get jobs
```

```
>>Kuberctl get pods
```

Tracks job completion status in real-time through Kubernetes watch mode.



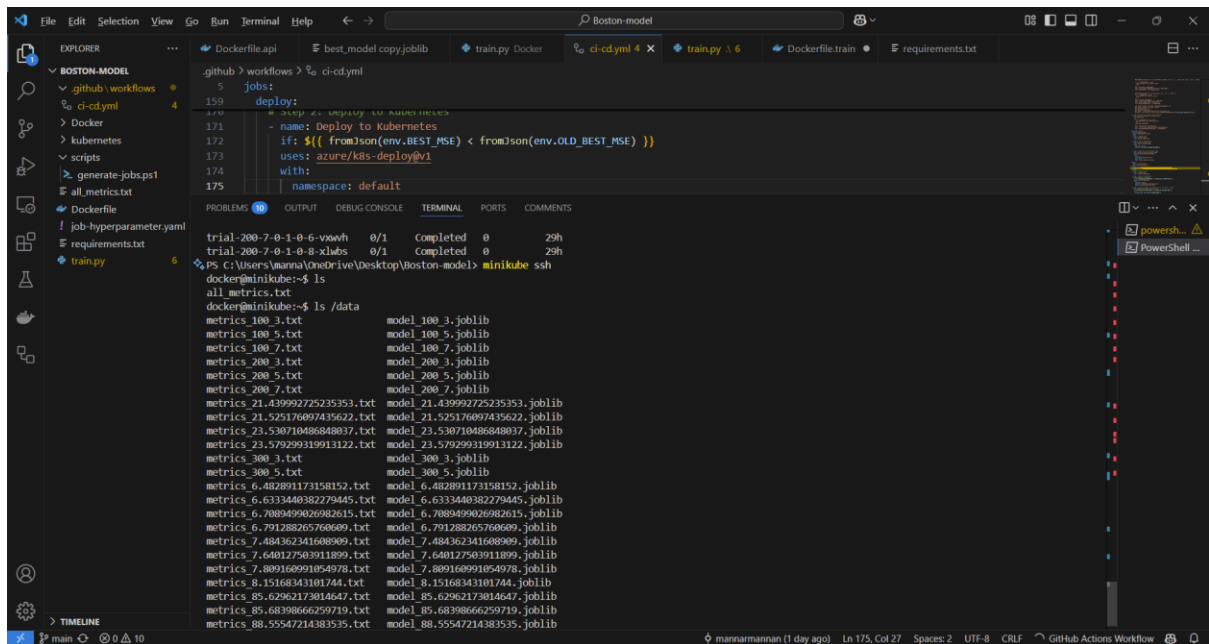
7.Login into shared memory of Virtual machine

```
>>Minikube ssh
```

Inspect Training Results

```
>>minikube ssh "ls -lh /data"
```

Verifies artifact creation in shared volume with model files and metrics.



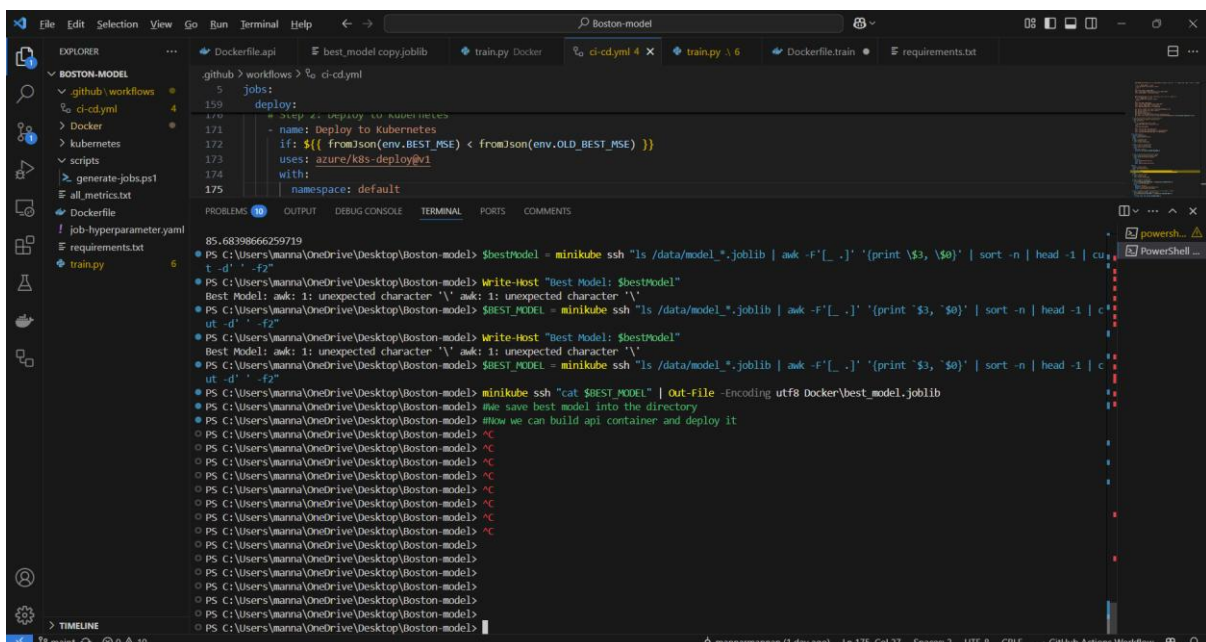
9. Retrieve Best Model

```
>>minikube ssh "ls /data/model *.joblib | sort -V | head -1"
```

```
>>$BEST_MODEL = minikube ssh "ls /data/model *.joblib | awk -F'[_]' '{print '$3, '$0'}' | sort -n | head -1 | cut -d' ' -f2"
```

```
>>minikube ssh "cat $BEST_MODEL" | Out-File -Encoding utf8 Docker\best_model.joblib
```

Uses version-sorting to identify the model with lowest MSE from filename patterns.



Task 2: Build a CI/CD Pipeline

CI/CD workflow file for all the above instruction to automate the process

CI/CD Workflow Enhancements:

- Added clearer stage separation (build → tune → deploy)
- Fixed environment variable handling for MSE comparisons
- Improved error handling in model retrieval steps
- Added explicit skip-ci tags to prevent pipeline loops

Key Pipeline Improvements:

1. Parallel Build Stages
Simultaneous image building and dependency installation for faster execution
2. Atomic Artifact Handling
Models and metrics stored with MSE in filenames for easy version comparison
3. Rolling Update Strategy
Conditional deployment only when new model outperforms previous best MSE
4. Persistent Metric Tracking
Commits best MSE to repository for historical comparison across runs

NOTE: Major efforts not taken in deployment stage

```
name: Model CI/CD Pipeline with Hyperparameter Tuning

on: [push]

jobs:
  build-and-train:
    runs-on: ubuntu-latest
    steps:
      # Step 1: Checkout code
      - name: Checkout code
        uses: actions/checkout@v2

      # Step 2: Set up Docker Buildx
      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v1

      # Step 3: Log in to Docker Hub
      - name: Log in to Docker Hub
        uses: docker/login-action@v1
      with:
```

```
username: mannarn
password: ${ secrets.DOCKER_PASSWORD }
```

Step 4: Build and push Train Docker image

```
- name: Build and push Train Docker image
  uses: docker/build-push-action@v2
  with:
    context: .
    file: Docker/Dockerfile.train
    push: true
    tags: mannarn/model-train:latest
```

hyperparameter-tuning:

```
runs-on: ubuntu-latest
needs: build-and-train
outputs:
  best_mse: ${ steps.collect_results.outputs.best_mse }
  best_model_file: ${ steps.collect_results.outputs.best_model_file }
  old_best_mse: ${ steps.retrieve_old.outputs.old_best_mse }
```

steps:

Step 1: Checkout code

```
- name: Checkout code
  uses: actions/checkout@v2
```

Step 2: Set up Python

```
- name: Set up Python
  uses: actions/setup-python@v2
  with:
    python-version: '3.8'
```

Step 3: Install dependencies

```
- name: Install dependencies
  run: pip install -r requirements.txt
```

Step 4: Set up Minikube

```
- name: Set up Minikube
  run: |
    curl -Lo minikube https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
    chmod +x minikube
    sudo mv minikube /usr/local/bin/
    minikube start --driver=docker --cpus=4 --memory=8192mb
```

Step 5: Generate Kubernetes Jobs for hyperparameter tuning

```
- name: Generate Kubernetes Jobs
  shell: pwsh
  run: |
    chmod +x scripts/generate-jobs.ps1
    ./scripts/generate-jobs.ps1
```

Step 6: Wait for Jobs to complete

- name: Wait for Jobs to complete

run: |

kubectl wait --for=condition=complete --timeout=1000s job --all || true

Step 7: Check Job and Pod Status

- name: Check Job and Pod Status

run: |

kubectl get jobs

kubectl get pods

Step 8: Retrieve logs of failed pods: Debugging

- name: Retrieve logs of failed pods

run: |

for pod in \$(kubectl get pods --field-selector=status.phase=Failed -o
jsonpath='{.items[*].metadata.name}'); do

echo "Logs for pod \$pod:"

kubectl logs \$pod

done

Step 9: Get the results of hyperparameter tuning

- name: Getting results

id: getting_results

run: |

BEST_MODEL=\$(minikube ssh "ls /data/model_*.joblib | awk -F'[_.]' '{print \\$3, \\$0}' | sort -n |
head -1 | cut -d' ' -f2")

if [-z "\$BEST_MODEL"]; then

echo "No model file found in /data/"

exit 1

fi

echo "Best Model: \$BEST_MODEL"

echo "::set-output name=best_model_file::\$BEST_MODEL"

echo "\$BEST_MODEL" > best_model.joblib

BEST_MSE=\$(minikube ssh "cat /data/mse_*.txt | sort -n | head -1")

if [-z "\$BEST_MSE"]; then

echo "No MSE file found in /data/"

exit 1

fi

echo "Best MSE: \$BEST_MSE"

echo "::set-output name=best_mse::\$BEST_MSE"

echo "Output: \$BEST_MSE" > best_mse.txt

echo "BEST_MSE=\$BEST_MSE" >> \$GITHUB_ENV

git config --global user.email "mannarmannan02@gmail.com"

git config --global user.name "mannarn"

```
git add best_mse.txt
git add best_model.joblib
git commit -m "Uploading model and mse scores[skip ci]"
git push --force https://x-access-token:${GITHUB_TOKEN}@github.com/${{ github.repository }}.git
```

Step 10: Retrieve old model's best MSE score

- name: Retrieve old model's best MSE score

id: retrieve_old

run: |

if [-f old_best_mse.txt]; then

OLD_BEST_MSE=\$(cat old_best_mse.txt)

else

OLD_BEST_MSE=999999

fi

echo "Old Best MSE: \$OLD_BEST_MSE"

echo "::set-output name=old_best_mse::\$OLD_BEST_MSE"

echo "OLD_BEST_MSE=\$OLD_BEST_MSE" >> \$GITHUB_ENV

build-and-push-api:

runs-on: ubuntu-latest

needs: hyperparameter-tuning

steps:

Step 1: Checkout code

- name: Checkout code

uses: actions/checkout@v2

Step 2: Log in to Docker Hub

- name: Log in to Docker Hub

uses: docker/login-action@v1

with:

username: mannarn

password: \${{ secrets.DOCKER_PASSWORD }}

Step 4: Build and push API Docker image

- name: Build and push API Docker image

uses: docker/build-push-action@v2

with:

context: .

file: Docker/Dockerfile.api

push: true

tags: mannarn/model-api:latest

deploy:

runs-on: ubuntu-latest

needs: build-and-push-api

env:

```

BEST_MSE: ${{ needs.hyperparameter-tuning.outputs.best_mse }}
OLD_BEST_MSE: ${{ needs.hyperparameter-tuning.outputs.old_best_mse }}
steps:
  # Step 1: Checkout code
  - name: Checkout code
    uses: actions/checkout@v2

  # Step 2: Deploy to Kubernetes
  - name: Deploy to Kubernetes
    if: ${{ fromJson(env.BEST_MSE) < fromJson(env.OLD_BEST_MSE) }}
    uses: azure/k8s-deploy@v1
    with:
      namespace: default
      manifests: kubernetes/deployment.yaml
      images: mannarn/model-api:latest

  # Step 3: Update old model's best MSE score if new model is better
  - name: Update old model's best MSE score
    if: ${{ fromJson(env.BEST_MSE) < fromJson(env.OLD_BEST_MSE) }}
    env:
      GITHUB_TOKEN: ${{ secrets.GITHUB_TOKEN }}
    run: |
      echo "$BEST_MSE" > old_best_mse.txt
      git config --global user.email "mannarmannan02@gmail.com"
      git config --global user.name "mannarn"
      git add old_best_mse.txt
      git commit -m "Update old best MSE score[skip ci]"
      git push --force https://x-access-token:${GITHUB_TOKEN}@github.com/${{ github.repository }}.git

```

Conclusion

In this project, I successfully developed and deployed an end-to-end machine learning pipeline for predicting house prices using the XGBoost regression model. The implementation involved hyperparameter tuning through Kubernetes Jobs, ensuring efficient parallel execution of multiple training trials.

By leveraging Kubernetes for parallelized hyperparameter tuning, I optimized model performance while maintaining scalability and automation. The integration of GitHub Actions enabled a robust CI/CD pipeline, automating the build, tuning, and deployment processes seamlessly.